

# MFE R Programming Workshop

Week 3

Brett Dunn and Mahyar Kargar

Fall 2017

# Introduction

# Questions

Any questions before we start?

# Overview of Week 3

- ▶ We will cover four classes to store array-like data:
  - ▶ Matrices
  - ▶ Data Frames
  - ▶ Tibbles
  - ▶ Intro to Data Tables

# Changing the Working Directory

- ▶ Use the `setwd()` R function
- ▶ Use the **Tools | Global Options | General** menu
- ▶ From within the Files pane, use the **More | Set As Working Directory** menu. (Navigation within the Files pane alone will not change the working directory.)

# Matrices vs Data Frames

- ▶ Matrix is a data type in R with the dimension attribute - the rows and the columns.
  - ▶ It has the elements of *same* class type.
  - ▶ We can have character, integer or complex elements in the matrices and so on.
  - ▶ We *cannot* have elements of mixed modes/class types such as both integer and character elements in the same matrix.
- ▶ A `data.frame` is a list of vectors of equal length, and the vectors can be of different types.
  - ▶ e.g. one character column, one numeric column.
- ▶ Tibbles and data tables inherit the functionality of data frames and improve on them in various ways.

# Matrices

# Creating Matrices

- ▶ Matrices are vectors with a number of rows and number of columns attribute.

```
myvec <- 1:10  
mymat <- matrix(myvec, nrow=2, ncol=5, byrow = FALSE)  
mymat
```

```
##           [,1] [,2] [,3] [,4] [,5]  
## [1,]        1    3    5    7    9  
## [2,]        2    4    6    8   10
```

```
dim(mymat)  # returns the dimension
```

```
## [1] 2 5
```



# Accessing Elements of Matrices

- ▶ Like vectors, elements can be accessed using `[]`

```
myamat <- matrix(1:15, nrow=3, ncol=5)
myamat[1, 2]  # row 1, column 2
```

```
## [1] 4
```

```
myamat[2:3, c(1, 4, 5)]
```

```
##      [,1] [,2] [,3]
## [1,]    2   11   14
## [2,]    3   12   15
```

## Filtering Matrices

- Filtering can be done on a single column or a single row, otherwise the filter returns a vector.

```
myvec <- c(1, 1, 3, 1, 5, 1, 7, 1, 9, 1)
mymat <- matrix(myvec, nrow=2, ncol=5)
mymat
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    3    5    7    9
## [2,]    1    1    1    1    1
```

```
mymat[, mymat[1, ] > 4]
```

```
##      [,1] [,2] [,3]
## [1,]    5    7    9
## [2,]    1    1    1
```

# Vectorization

- Most R functions work on matrices as well.

```
mymat <- matrix(1:10, nrow=2, ncol=5)  
exp(mymat)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]  
## [1,]  2.718282 20.08554 148.4132 1096.633  8103.084  
## [2,]  7.389056 54.59815 403.4288 2980.958 22026.466
```

```
sd(mymat)  # standard deviation
```

```
## [1] 3.02765
```

# Applying Functions to Rows and Columns

- ▶ `apply` allows you to apply a function across a dimension of a matrix.
- ▶ The third argument is a function!

```
mymat <- matrix(1:10, nrow=2)  
# mean across rows (can also use rowMeans() function)  
apply(mymat, 1, mean) # apply mean along rows
```

```
## [1] 5 6
```

```
apply(mymat, 2, max) # apply max along columns
```

```
## [1] 2 4 6 8 10
```

# Combining Matrices with cbind and rbind

- ▶ Column bind and row bind.

```
mymat1 <- matrix(1:4, nrow=2)
mymat2 <- matrix(6:9, nrow=2)
mymat3 <- matrix(10:11, ncol=2)
cbind(mymat1, mymat2)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    3    6    8
## [2,]    2    4    7    9
```

```
rbind(mymat1, mymat3)
```

```
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
## [3,]   10   11
```

# Matrix operations

- ▶ Many matrix operations are surrounded by % signs.

```
mymat1 <- matrix(1:4, nrow=2)
mymat2 <- matrix(5:8, nrow=2)
mymat1 %*% mymat2  # matrix multiplication
```

```
##      [,1] [,2]
## [1,]   23  31
## [2,]   34  46
```

```
mymat1 + mymat2
```

```
##      [,1] [,2]
## [1,]    6  10
## [2,]    8  12
```

# Matrix Algebra

- ▶ <http://www.statmethods.net/advstats/matrix.html>

# Data Frames



## data.frames

- ▶ The `data.frame` is one of the most useful features in R.
- ▶ A `data.frame` is like a `matrix` with a two-dimensional rows-and-columns structure.
- ▶ However, unlike a `matrix`, in a `data.frame` each column can have a *different* data type.
  - ▶ For example, one column might be numbers and another characters.
- ▶ Technically, a `data.frame` is a `list`, with the components of the list being equal-length vectors.
- ▶ Each column must be the same length (unlike a `list`).

## Creating data.frames

- ▶ Unless you are working with categorical data, you probably want to set `stringsAsFactors=FALSE`.

```
courses <- c("Stochastic Calculus", "Fixed Income")
examGrades <- c(92, 98)
gradeBook <- data.frame(courses, examGrades,
                        stringsAsFactors = FALSE)
gradeBook
```

```
##           courses examGrades
## 1 Stochastic Calculus      92
## 2      Fixed Income      98
```

# Column Names

- ▶ Column names in `data.frames` are specified by `names()`.
- ▶ This is because `data.frames` are actually lists with special attributes.
- ▶ That means that the usual list functions work on `data.frames`.
- ▶ `lapply`, etc.

## Accessing Elements of `data.frames`

## Accessing Elements of data.frames

- ▶ We can access a data.frame component just like a list.

```
gradeBook[[1]]  # first way
```

```
## [1] "Stochastic Calculus" "Fixed Income"
```

```
gradeBook[["courses"]]  # second way
```

```
## [1] "Stochastic Calculus" "Fixed Income"
```

```
gradeBook$courses  # third way
```

```
## [1] "Stochastic Calculus" "Fixed Income"
```

## Accessing Elements of data.frames (2)

- Note that `[[` simplifies the result, returning a vector:

```
str(gradeBook[[1]])
```

```
## chr [1:2] "Stochastic Calculus" "Fixed Income"
```

- But `[` is preserving subsetting and (usually) returns a 'data.frame':

```
str(gradeBook[1])
```

```
## 'data.frame': 2 obs. of 1 variable:
```

```
## $ courses: chr "Stochastic Calculus" "Fixed Income"
```

## Accessing Elements of data.frames like a matrix

- We can access data.frame elements like a matrix.

```
gradeBook[1,2]
```

```
## [1] 92
```

```
gradeBook[1,] # the first row
```

```
##           courses examGrades  
## 1 Stochastic Calculus          92
```

```
gradeBook[,2] # the second column
```

```
## [1] 92 98
```

## Accessing Elements of data.frames like a matrix(2)

- ▶ To preserve the data.frame class, set drop = FALSE

```
gradeBook[,2] # returns a vector
```

```
## [1] 92 98
```

```
gradeBook[,2,drop=FALSE] # returns a data.frame
```

```
##    examGrades
```

```
## 1          92
```

```
## 2          98
```



## Subsetting data.frames

## Filtering with subset()

```
set.seed(1234)
x.df <- data.frame(V1 = rnorm(4), V2 = runif(4),
                   V3 = rchisq(4, df = 2), V4 = 1:4)
x.df
```

##		V1	V2	V3	V4
## 1	-1.2070657	0.6660838	0.3523580	1	
## 2	0.2774292	0.5142511	2.8742845	2	
## 3	1.0844412	0.6935913	0.3134394	3	
## 4	-2.3456977	0.5449748	0.4390040	4	

```
x.sub <- subset(x.df, V4 > 2)
x.sub
```

##		V1	V2	V3	V4
## 3	1.084441	0.6935913	0.3134394	3	
## 4	-2.345698	0.5449748	0.4390040	4	

## Subsetting rows using conditional statements

- ▶ The data frame `x.sub1` contains only the observations for which the values of the variable `V4` is greater than 2 and the variable `V1` is greater than 0.6.

```
x.sub1 <- subset(x.df, V4 > 2 & V1 > 0.6)
x.sub1
```

```
##           V1           V2           V3 V4
## 3 1.084441 0.6935913 0.3134394 3
```

## Subsetting both rows and columns

- ▶ The data frame `x.sub2` contains only the variables `V2` and `V3` and then only the observations of these two variables where the values of variable `V4` are greater than 2 and the values of variable `V1` are greater than 0.6.

```
x.sub2 <- subset(x.df, V4 > 2 & V1 > 0.6,  
                 select = c(V2, V3))  
x.sub2
```

```
##           V2           V3  
## 3 0.6935913 0.3134394
```

## Subsetting rows using indices

- ▶ The `x.sub3` data frame contains only the observations for which the values of variable `V4` are equal to 2.

```
x.sub3 <- x.df[x.df$V4 == 2, ]  
x.sub3
```

```
##           V1           V2           V3 V4  
## 2 0.2774292 0.5142511 2.874284  2
```

## Subsetting rows using %in%

- ▶ The x.sub4 data frame contains only the observations for which the values of variable V4 are equal to either 1 or 4.

```
x.sub4 <- x.df[x.df$V4 %in% c(1, 4), ]  
x.sub4
```

##		V1	V2	V3	V4
##	1	-1.207066	0.6660838	0.352358	1
##	4	-2.345698	0.5449748	0.439004	4

## Subsetting columns using indices

- ▶ The `x.sub5` data frame contains all the rows on `x.df`, removing the first and third columns

```
x.sub5 <- x.df[, -c(1,3)]  
x.sub5
```

```
##           V2 V4  
## 1 0.6660838  1  
## 2 0.5142511  2  
## 3 0.6935913  3  
## 4 0.5449748  4
```

## Complete.cases

- ▶ `complete.cases()` gets rid of any rows with at least one NA value.

```
# Let's makes the second col in row 1 an NA  
x.df[1,2] <- NA  
x.df[complete.cases(x.df), ] # removes row 1
```

##		V1	V2	V3	V4
## 2		0.2774292	0.5142511	2.8742845	2
## 3		1.0844412	0.6935913	0.3134394	3
## 4		-2.3456977	0.5449748	0.4390040	4



## Merging data.frames

- ▶ Two data.frames can be combined using the merge function.

```
courses <- c("Stochastic Calculus", "Fixed Income")
midtermGrades <- c(89, 91)
gradeBook2 <- data.frame(courses, midtermGrades,
                          stringsAsFactors = FALSE)
merge(gradeBook, gradeBook2)
```

	courses	examGrades	midtermGrades
## 1	Fixed Income	98	91
## 2	Stochastic Calculus	92	89

## Adding Columns to data.frames

```
dat1 <- 1:4
dat2 <- rep(c("A","B"),each=2)
myframe <- data.frame(col1=dat1,col2=dat2)
myframe$col3 <- 5:8
myframe
```

##	col1	col2	col3
## 1	1	A	5
## 2	2	A	6
## 3	3	B	7
## 4	4	B	8

## Reading in Data from a CSV File

- ▶ Reading in data typically gives you a `data.frame`.
- ▶ `read.table` is the basic function to read in tabular data.
- ▶ `read.csv` is a special case of `read.table`.
- ▶ As usual see `?read.table`.
- ▶ Often you want to set `stringsAsFactors = FALSE`.
- ▶ `write.csv` writes data to a `.csv` file.

```
optdata <- read.csv(file="../week2/lab/optionsdata.csv",  
                    header = T, stringsAsFactors = FALSE)  
head(optdata, 3)
```

```
##      S0 sigma    r T    K  
## 1 100    0.3 0.0 1 100  
## 2 101    0.3 0.0 1 100  
## 3 101    0.1 0.1 1 105
```

tibbles

# What is a tibble?

- ▶ Tibbles are a “modern take” on R’s traditional `data.frame`.
- ▶ They keep the features that have stood the test of time, and drop the features that used to be convenient but are now frustrating (i.e. converting character vectors to factors).

## Creating a tibble

- ▶ `tibble()` can be used to create a data frame.
- ▶ It never changes an input's type (i.e., no more `stringsAsFactors = FALSE!`).

```
tibble(x = letters)
```

```
## # A tibble: 26 × 1
```

```
##       x
```

```
##   <chr>
```

```
## 1     a
```

```
## 2     b
```

```
## 3     c
```

```
## 4     d
```

```
## 5     e
```

```
## 6     f
```

```
## 7     g
```

```
## 8     h
```

```
## 9     i
```

## Creating a tibble of lists

This makes it easier to use with list-columns:

```
tibble(x = 1:3, y = list(1:5, 1:10, 1:20))
```

```
## # A tibble: 3 × 2
##       x         y
##   <int>   <list>
## 1     1 <int [5]>
## 2     2 <int [10]>
## 3     3 <int [20]>
```

# Lazy and Sequential Evaluation

It evaluates its arguments lazily and sequentially:

```
tibble(x = 1:5, y = 1, z = x ^ 2 + y)
```

```
## # A tibble: 5 × 3
##       x     y     z
##   <int> <dbl> <dbl>
## 1     1     1     2
## 2     2     1     5
## 3     3     1    10
## 4     4     1    17
## 5     5     1    26
```



# Column Names

- ▶ Tibbles never adjust the names of variables:

```
names(data.frame(`crazy name` = 1))
```

```
## [1] "crazy.name"
```

```
names(tibble(`crazy name` = 1))
```

```
## [1] "crazy name"
```

## Other Features

- ▶ Tibbles don't use `row.names()`. It never stores a variable as special attribute.
- ▶ Tibbles only recycle vectors of length 1. This is because recycling vectors of greater lengths is a frequent source of bugs.
- ▶ Tibble provides `as_tibble()` to coerce objects into tibbles.

# Tibbles vs Data Frames: Printing

- ▶ When you print a tibble, it only shows the first ten rows and all the columns that fit on one screen. It also prints an abbreviated description of the column type.
- ▶ You can control the default appearance with options.

```
options(tibble.print_max = 3, tibble.print_min = 2)
tibble(x = 1:1000)
```

```
## # A tibble: 1,000 × 1
##       x
##   <int>
## 1     1
## 2     2
## # ... with 998 more rows
```

## Tibbles vs Data Frames: Subsetting

- ▶ If you want to pull out a single variable, you could use \$ and [[ or %>% ([[ can extract by name or position; \$ only extracts by name)

```
df <- tibble( x = 1:5, y = rnorm(5))  
# Extract by name (df[["x"]] works the same way)  
df$x
```

```
## [1] 1 2 3 4 5
```

```
# Extract by position  
df[[1]]
```

```
## [1] 1 2 3 4 5
```

```
df %>% .$x # df %>% .[["x"]] works the same way
```

```
## [1] 1 2 3 4 5
```

## Tibbles vs Data Frames: Subsetting with `[[` and `$`

- ▶ Recall that for data frames:
  - ▶ `[[` extracts a single column as a vector.
  - ▶ `$` works similarly to `[[`, but does *partial matching* on the column name.

```
df <- data.frame(colName = 1:5, m = 2:6); df$c
```

```
## [1] 1 2 3 4 5
```

- ▶ Tibbles never do partial matching, and will throw a warning and return NULL if the column does not exist.

```
tbl <- as_tibble(df); tbl$c
```

```
## Warning: Unknown or uninitialised column: 'c'.
```

```
## NULL
```

# Tibbles vs Data Frames: Recycling

- ▶ When constructing a tibble, only values of length 1 are recycled.
- ▶ The first column with length different to one determines the number of rows in the tibble, conflicts lead to an error.
  - ▶ `tibble(a = 1:3, c = 1:2)` gives “Error: Column c must be length 1 or 3, not 2”.

Very quick intro to data.tables

# What is a `data.table`?

- ▶ Think of `data.table` as an advanced version of `data.frame`.
  - ▶ Every column is the same length, but may have a different type
- ▶ It inherits from `data.frame` and works even when `data.frame` syntax is applied on `data.table`
- ▶ `data.table` is *very fast*.
- ▶ It is one of the most useful packages in R.
- ▶ The syntax of `data.table` is concise.
  - ▶ Lowers programmer time...
  - ▶ ...but it can be hard to understand
  - ▶ Make sure you comment your code!

```
library(data.table)
```



## An Example

- ▶ Syntax is DT[i, j, by]:
- ▶ “Take DT, subset rows using i, then calculate j grouped by by.”

```
data("mtcars")
mtcarsDT <- data.table(mtcars)
mtcarsDT[ mpg > 20,
          .(AvgHP = mean(hp), "MinWT(kg)" = min(wt*453.6)),
          by = .(cyl, under5gears = gear < 5)]
```

	cyl	under5gears	AvgHP	MinWT(kg)
## 1:	6	TRUE	110.00000	1188.4320
## 2:	4	TRUE	78.33333	732.5640
## 3:	4	FALSE	102.00000	686.2968

# Why learn `data.table`?

- ▶ For data that fits in memory, `data.table` is much faster than `data.frames` and `tibbles`.
- ▶ `data.table` also saves memory because it avoids copying large objects.
- ▶ Part of the speed advantage comes from the fact that `data.table` provides a set of tools to update by reference.
  - ▶ In base R, if a function modifies a single element of a large `data.frame`, a copy of the entire `data.frame` is made.
- ▶ `data.table` provides a powerful set of commands to access data.
- ▶ We will cover `data.table` in much more detail later in the course.